

— AI4I TRACK 3 — DEVELOPMENT

NeuroWorks Conditional Advance

Evaluator scorecard for the 2026 AI for Impact Challenge. A *build-led* entry with an impact-evidence gap.

PREPARED FOR

OrionX Evaluation Panel

DATE

13 July 2026

COUNTERPARTY

Neuroworks / RUBIEM Developers

FORMAT

A4 • 920px

01

— EVALUATION SUMMARY

Advance — Conditional

Score 69/100. Shortlist for the Mutare bootcamp on the strength of a rare, verifiable, working build — but condition advancement on three diligence items. This is a build-led entry with an impact-evidence gap, not an idea-led entry with an execution gap.

OUTCOME

Shortlist for the Mutare bootcamp (27 Jul – 1 Aug). Advancement conditional on verified impact evidence, corrected NDS2 mapping, and full team disclosure.

TOTAL SCORE

69/100

Top of the conditional band

TECHNICAL MERIT

18/20

Strongest dimension

TESTS PASSING

118

Live test suite

LOCAL INFERENCE

62.7 %

Calls on consumer RTX 3060

RESPONSIBLE AI

12/15

Governance-first posture

LIVE ROUTERS

28

Production endpoints

RUBRIC NOTE

The official AI4I Terms of Reference rubric was not accessible at evaluation time. Scores use a transparent 100-point model reconstructed from the challenge's stated priorities (impact-first, responsible-AI/governance-heavy, sustainability, national-strategy alignment). Swap in official weights if available — the per-criterion rationale holds regardless of weighting.

CRITERION	WT.	SCORE	RATIONALE
Impact & Strategic Alignment	25	16	Latent alignment is high but under-claimed and impact is undemonstrated. Correctly cites AI Strategy Pillars 1 & 3; misses Pillars 2 (computational sovereignty) and 4 (governance & ethics) it clearly satisfies. NDS2 mapping appears factually wrong. Zero signed pilots — impact is asserted via capability, not measured.
Technical Merit & Feasibility	20	18	Strongest dimension. A working, verifiable system: 28 live routers, ~90 primitives, 118 passing tests, live PostgreSQL/Paynow/Mailjet/Stripe, root-hashed datasets, checkable endpoints. Feasibility is proven, not promised.
Innovation & Originality	15	10	Low as AI research (agentic workforce, HITL tool-calling, RAG, local inference are 2025 best-practice, crowded globally). High as a localized, governed sovereign system. IntelliNexus governed-knowledge flywheel is the one genuinely novel thread — underdeveloped.
Responsible AI, Governance & Data Protection	15	12	Governance-first, not bolt-on: pre-execution policy enforcement, mandatory HITL, full audit trail, AES-256-GCM, DPA controller/processor analysis, radical candor incl. disclosed-and-rotated credential leak. Deductions: bias unevaluated; Shona/Ndebele appropriateness not linguistically reviewed.
Sustainability & Scalability	15	8	Architecture scales cheaply (local-first, measured \$0.10/task). But the business model is asserted, not modeled: break-even org counts with no pricing, CAC, or conversion path; no LOIs; managed relay not built.
Team & Execution Capability	10	5	Single named lead (Aiia Rubiem / RUBIEM Developers) against a very large scope. No team section, roles, or capacity evidence. Bus-factor and roadmap-execution risk unaddressed in a team-based challenge.
TOTAL	100	69	Shortlist-worthy on execution; conditional on impact evidence & team.

02

— DIFFERENTIATORS

What Sets It Apart

Four decisive strengths separate this entry from the typical pitch deck.

*It runs and is checkable. A working system with verifiable artifacts in a category dominated by **decks**.*

- **It runs and is checkable.** A working system with verifiable artifacts — endpoints, root hashes, passing tests — in a track dominated by slide decks. Decisive for a Development-track entry.
- **Sovereignty by default.** 62.7% of calls served locally on a consumer RTX 3060; data never leaves the machine by default — directly on-strategy for a forex- and connectivity-constrained sovereign context.
- **Governance-first posture and candor.** Pre-execution policy enforcement plus mandatory HITL plus audit trail, plus honest disclosure of gaps and a real, disclosed-and-rotated security incident. Rare and trust-building.
- **Deep local wiring.** Paynow rail already live; Shona and Ndebele support; sector packs named to real institutions: PRAZ, MOHCC, ZIMSEC, GMB, ZTA, ZIMRA.

03

— RISK ASSESSMENT

Red Flags to Probe

Five gaps that could block final funding if not closed at the bootcamp.

- **No demonstrated impact.** Zero signed pilots; one live database that appears to be their own. For an "AI for Impact" award, this is the load-bearing weakness.
- **Probable factual error on NDS2.** Presents a 5-pillar NDS2 with invented names; NDS2 (2026–2030) is built on ~10 national priorities. Asserting this to the strategy's own authors is a credibility risk.
- **Team opacity.** One named individual against a very large surface area — either a small team carrying roadmap risk, or an implausible scope-to-team ratio. Unaddressed.
- **Sustainability asserted, not modeled.** Break-even org counts with no pricing, CAC, or conversion path. No LOIs. Managed relay not built.
- **Adoption vs. complexity mismatch.** Claims non-technical operators; the evidence surface (37-page console, HERMES allowlist, npm/repo verification) is developer-centric.

04

— BOOTCAMP DILIGENCE

Questions for Mutare, 27 Jul – 1 Aug

Eight concrete items to resolve during the 27 July – 1 August bootcamp.

- 01 Show one end-to-end real-world task on a real Zimbabwean organisation's data with a before/after metric (hours saved, error rate). Not internal operational data — an external workflow.
- 02 Do you have any signed LOI/MOU with a named institution? If not, what is the concrete path to first signed pilot, and what is blocking it today?
- 03 Who is the full team? Names, roles, and who owns which of the 28 routers / 90 primitives / data pipeline. What is the bus factor?
- 04 Reconcile your NDS2 pillar claims against the actual NDS2 document. Which national priorities do you actually map to?
- 05 Walk a genuinely non-technical operator (e.g. a rural clinic administrator) through first mission setup on screen. How long, how many decisions?
- 06 37.3% of measured calls went to paid cloud. What classes of task fail on local models, and what is the real (not best-case) cost at 20-org scale?
- 07 What is your bias / linguistic-appropriateness evaluation plan for Shona/Ndebele, and who is the named linguistic authority?
- 08 The platform code is proprietary while the pitch is national data infrastructure. What prevents vendor lock-in for a sovereign government deployment?

OUTCOME

Closing the impact-evidence gap at the bootcamp would move this entry into the strong band (70–84).

Scoring Scale & Framework

Weights reconstructed from published AI4I priorities; not the official ToR rubric. Strategy-pillar and NDS2 findings verified against the Zimbabwe National AI Strategy 2026–2030 and NDS2 2026–2030.

BAND	RANGE	MEANING
Priority fund	85–100	Immediate funding recommendation
Strong, advance	70–84	Clear bootcamp placement
Advance, conditional	55–69	Bootcamp with material conditions
Borderline	40–54	Unlikely without major rewrite
Decline	<40	Does not meet threshold

NeuroWorks scores 69 — top of the conditional band, adjacent to strong. The entry is shortlist-worthy on execution. Closing the impact-evidence gap at the bootcamp would move it into the strong band.

VERIFICATION

Strategy-pillar and NDS2 findings verified against the Zimbabwe National AI Strategy 2026–2030 and NDS2 2026–2030.

